

Machine Learning in Databricks with Spark ML

Brandon Ritchie, Rachelle Ceron, Spencer Carillo, Gavin South, Trevor Neri

Spark MLlib (Dataframe Based) vs Spark MLlib (RDD Based)

MLlib (RDD)

- Deprecated spark machine learning library
- Built on RDD's (<https://spark.apache.org/docs/latest/ml-guide.html>)

MLlib (DataFrame)

- Built on dataframes that translate to the RDD structure
- More user friendly API.

Important Note

Both methods are actually inefficient with smaller datasets. If you are working with smaller data, sci-kit learn is the way to go with RAM stored data.

What is an RDD?

Resilient Distributed Dataset - A collection of elements of data, partitioned across nodes that allows processes to be run in parallel.

The dataframe structure in Spark ML allows for seamless translation to RDD's.

(<https://databricks.com/glossary/what-is-rdd>)

5 Minute Machine Learning Review Dump

Preprocess Data - Manipulating data to enhance performance of model

This link will be your best friend!!! (<https://spark.apache.org/docs/1.4.1/ml-features.html>)

- Feature engineering (done)
- Quantify categorical data
 - One-hot encoding
 - <https://stackoverflow.com/questions/42295001/how-to-interpret-results-of-spark-onehotencoder>
 - Label encoding
 - Clustering
- Interpolation to fill NA's
 - <https://www.analyticsvidhya.com/blog/2021/06/power-of-interpolation-in-python-to-fill-missing-values/>
- Split data into training and testing
 - <https://spark.apache.org/docs/2.1.0/ml-tuning.html#train-validation-split>
- Feature scaling (<https://towardsdatascience.com/normalization-vs-standardization-quantitative-analysis-a91e8a79cebf>)
 - Prevents large features from dominating the model
 - Standardization - Centers data with mean 0 and a standard deviation of 1
 - Normalization - Rescales values from 0-1 -Spark ML StandardScalar()

5 Minute Machine Learning Review Dump

Fit Model to Data

Refer to documentation for Spark ML models under MLlib (DataFrame-based) tab

<https://spark.apache.org/docs/latest/api/python/reference/pyspark.ml.html>

5 Minute Machine Learning Review Dump

Evaluate Model

- Accuracy
 - Am I predicting more correctly than incorrectly
 - $(TP + TN) / (TP + TN + FP + FN)$
- Precision
 - How well can you pick 5 bad apples out of 1,000,000 total apples
 - $TP / (TP + FP)$
- Recall
 - How many bad apples are predicted out of all of the bad apples
 - TP / FN
- F-Score (No strong goal and want to do well with both precision and recall)
 - Harmonic mean of precision and recall
 - Less interpretable, but big number good small number bad
- Validation Curves
 - Plot error by training iteration for training and validation data

Coding Walkthrough

Follow the link below and make sure to clone it to your repository!

[Put link to walkthrough]